

# IT353 : Deep Learning

## Lab Assignment 3 - Audio and Text Classification

**Kesanam Ashinee**

211AI023

### Dataset 1: (For Audio Classification)

The dataset describes 440 recordings of cat meows categorized into three different situations:

- Waiting for food: This category likely contains recordings of meows associated with hunger or anticipation of receiving food.
- Isolation in unfamiliar environment: This category likely includes meows associated with stress or anxiety from being in an unfamiliar place.
- Brushing (being brushed affectionately by the owner): This category likely contains recordings of meows associated with positive interactions or contentment during petting.

Additional details:

- The dataset involves 21 cats belonging to two breeds: Maine Coon and European Shorthair.
- The purpose of the dataset is potentially to contribute to understanding cat vocalizations and their correlation with specific contexts or emotions.

Link: [Cat Meow Classification Dataset](#)

## Dataset 2: (For Text Classification)

This dataset is designed for classifying e-commerce product descriptions into four categories:

- Electronics: Products related to electronic devices and accessories.
- Household: Products related to household items and appliances.
- Books: Products related to books and publications.
- Clothing & Accessories: Products related to clothing, footwear, and accessories.

The dataset contains two columns:

- Class Name: This column specifies the category of the product description, represented as a string ("Electronics", "Household", "Books", or "Clothing & Accessories").
- Product Description: This column contains the textual description of the product, providing information about its features, specifications, and other relevant details.

Dataset Link: [E-Commerce Text Classification Dataset](#)

## Kaggle Notebook

All code can be found here - [Task 1 Notebook](#) , [Task 2 Notebook](#)

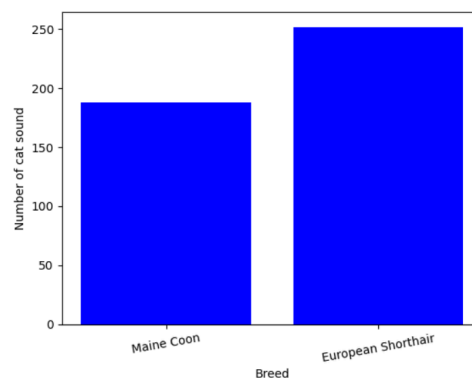
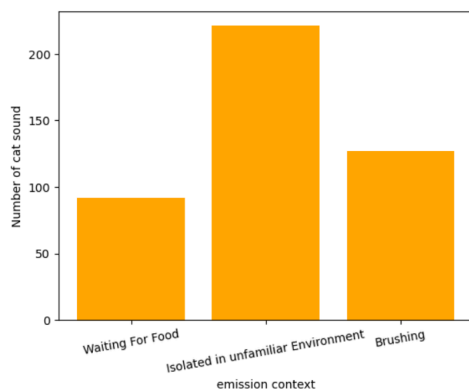
## Task 1:

Download any Audio Classification dataset (Birds sound, Gender, vehicle sound etc) and convert it into spectrogram, use Deep CNN models and classify the spectrograms.

## Implementation:


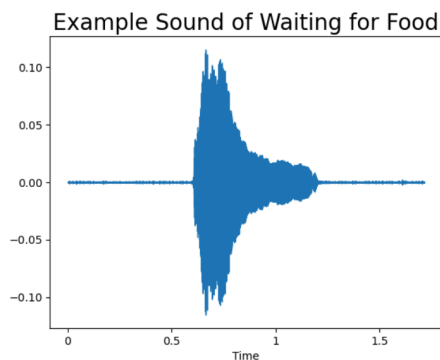
EDA:

Checking the value counts for each category





Data Loading and Visualization:


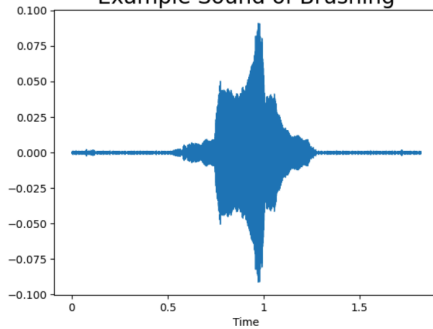
- The code loads audio files using librosa and visualizes the sound waveforms and spectrograms.
- The 3 types of meow are loaded and displayed.



Example Sound of Isolated in unfamiliar Environment



Example Sound of Brushing



## Data Augmentation Functions:

- Functions are defined for data augmentation:
  - Adding noise to the audio data.
  - Stretching the audio data in time.
  - Pitch shifting the audio data.

## Feature Extraction:

- Features are extracted from the audio data to be used for classification.
- Features include Zero Crossing Rate (ZCR), Chroma\_stft, Mel-frequency cepstral coefficients (MFCC), Root Mean Square (RMS) value, and MelSpectrogram.

## Feature Extraction Function:

- The `extract_features()` function computes various features from the audio data.

## Feature Extraction from Augmented Data:

- Features are extracted from original data, noise-added data, and stretched/pitched data.

# Model 1

## Model Architecture:

- The model is a sequential model defined using Keras's Sequential API.
- It comprises several layers including Conv1D, MaxPooling1D, Flatten, and Dense layers.

Model: "sequential"		
Layer (type)	Output Shape	Param #
=====		
conv1d (Conv1D)	(None, 160, 128)	512
max_pooling1d (MaxPooling1D)	(None, 160, 128)	0
conv1d_1 (Conv1D)	(None, 158, 256)	98560
max_pooling1d_1 (MaxPooling1D)	(None, 158, 256)	0
conv1d_2 (Conv1D)	(None, 156, 512)	393728
max_pooling1d_2 (MaxPooling1D)	(None, 156, 512)	0
conv1d_3 (Conv1D)	(None, 154, 1024)	1573888
max_pooling1d_3 (MaxPooling1D)	(None, 154, 1024)	0
flatten (Flatten)	(None, 157696)	0
dense (Dense)	(None, 512)	80740864
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131328
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 3)	387
=====		
Total params: 82,972,163		
Trainable params: 82,972,163		
Non-trainable params: 0		

## Convolutional Layers:

- Four Conv1D layers with increasing number of filters (128, 256, 512, and 1024 respectively).
- Each Conv1D layer uses a kernel size of 3 and ReLU activation function.
- Convolutional layers are designed to capture hierarchical patterns in the input data.

## MaxPooling Layers:

- Four MaxPooling1D layers with pool size (1).
- MaxPooling layers reduce the spatial dimensions of the input, helping in feature extraction and reducing computational complexity.

#### Flatten Layer:

- After the convolutional layers, there's a Flatten layer to flatten the 1D output from the convolutional layers into a 1D vector.

#### Dense Layers:

- Three Dense layers follow the Flatten layer with decreasing number of units (512, 256, and 128 respectively).
- Each Dense layer uses ReLU activation function.
- Dropout layers with a dropout rate of 0.3 are added after each Dense layer to prevent overfitting by randomly dropping a fraction of the input units.

#### Output Layer:

- The final Dense layer consists of 3 units with softmax activation, representing the probabilities of the input belonging to each of the three classes ('Waiting For Food', 'Isolated in unfamiliar Environment', 'Brushing').

#### Model Compilation:

- The model is compiled using categorical cross-entropy loss function, RMSprop optimizer, and accuracy as the evaluation metric.


#### Model Training:

- The model is trained for 30 epochs with a batch size of 128.
- Training data and validation data are passed during the training process.

#### Model Evaluation:


- After training, the model's accuracy is evaluated on the test data.
- Training and testing loss and accuracy are plotted over epochs to visualize the model's performance.

Accuracy of our model on test data : 67.42424368858337 %



### Confusion Matrix and Classification Report:

- Confusion matrix and classification report are generated to evaluate the model's performance on the test data.
- The confusion matrix helps in understanding the model's prediction across different classes.
- The classification report provides precision, recall, and F1-score for each class.



## Classification Report

	precision	recall	f1-score	support
0	0.77	0.40	0.52	91
1	0.67	0.89	0.77	123
2	0.60	0.64	0.62	50
accuracy			0.67	264
macro avg	0.68	0.64	0.64	264
weighted avg	0.69	0.67	0.65	264

## Model 2:

Model Architecture 2:


Similar to Model 1, some parameters are changed

Model: "sequential_1"		
Layer (type)	Output Shape	Param #
=====		
conv1d_4 (Conv1D)	(None, 160, 256)	1024
max_pooling1d_4 (MaxPooling 1D)	(None, 160, 256)	0
conv1d_5 (Conv1D)	(None, 158, 128)	98432
max_pooling1d_5 (MaxPooling 1D)	(None, 158, 128)	0
conv1d_6 (Conv1D)	(None, 156, 128)	49280
max_pooling1d_6 (MaxPooling 1D)	(None, 156, 128)	0
conv1d_7 (Conv1D)	(None, 154, 1024)	394240
max_pooling1d_7 (MaxPooling 1D)	(None, 154, 1024)	0
flatten_1 (Flatten)	(None, 157696)	0
dense_4 (Dense)	(None, 1024)	161481728
dropout_3 (Dropout)	(None, 1024)	0
dense_5 (Dense)	(None, 256)	262400
dropout_4 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 128)	32896
dropout_5 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 3)	387
=====		
Total params: 162,320,387		
Trainable params: 162,320,387		
Non-trainable params: 0		

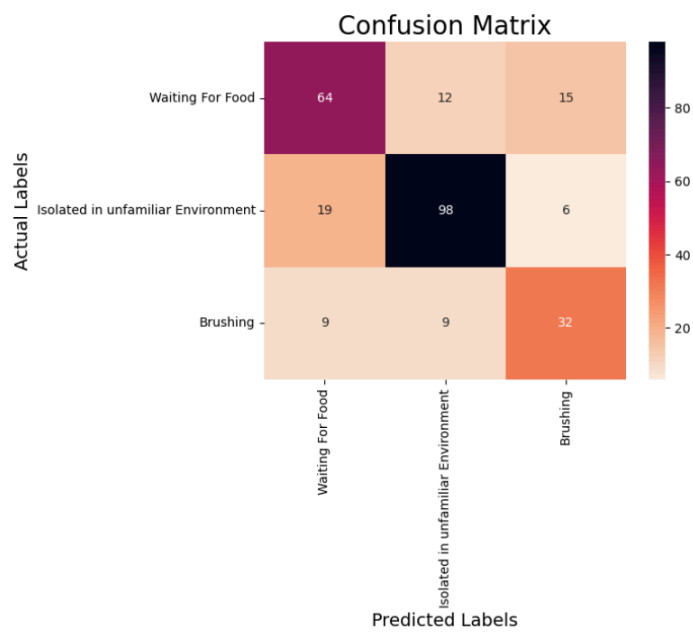


The results of Model 2 are follows:

Accuracy of our model on test data : 73.48484992980957 %



Confusion Matrix



Classification Report

	precision	recall	f1-score	support
0	0.70	0.70	0.70	91
1	0.82	0.80	0.81	123
2	0.60	0.64	0.62	50
accuracy			0.73	264
macro avg	0.71	0.71	0.71	264
weighted avg	0.74	0.73	0.74	264


## Task 2:

Download any Text Classification dataset (social media text or any short text) and convert it into meaningful image form!!, use Deep CNN models and classify the images.

## Implementation:

### Data Preprocessing:

- Calculated the value counts of the target variable




- The dataset contains textual data that needs to be converted into numerical format for machine learning models to process.
- BERT (Bidirectional Encoder Representations from Transformers) is used to convert the text data into word embeddings, which capture the semantic meaning of words in a high-dimensional space.

### BERT Integration:

- The BERT model and tokenizer are imported from the Hugging Face Transformers library.
- BERT is used to generate word embeddings for each sentence in the dataset.
- The output from BERT is processed to extract the word embeddings.
- It is visualized with the text embeddings. The embeddings are multiplied with 255 and is converted into a grey scale image

## Text Embedding in Black and White Image Representation



### Convolutional Neural Network (CNN) Model:

- Two CNN models are implemented to classify the text data.
- Model 1:

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
conv1d_7 (Conv1D)	(None, 766, 128)	195,968
max_pooling1d_7 (MaxPooling1D)	(None, 383, 128)	0
flatten_7 (Flatten)	(None, 49024)	0
dense_14 (Dense)	(None, 64)	3,137,600
dense_15 (Dense)	(None, 10)	650

Total params: 3,334,218 (12.72 MB)

Trainable params: 3,334,218 (12.72 MB)

Non-trainable params: 0 (0.00 B)

- Model 2:

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
conv1d_5 (Conv1D)	(None, 383, 256)	522,496
max_pooling1d_5 (MaxPooling1D)	(None, 127, 256)	0
flatten_5 (Flatten)	(None, 32512)	0
dense_10 (Dense)	(None, 64)	2,080,832
dense_11 (Dense)	(None, 4)	260

Total params: 2,603,588 (9.93 MB)

Trainable params: 2,603,588 (9.93 MB)

Non-trainable params: 0 (0.00 B)

- The first CNN model (model1) comprises a single Conv1D layer followed by MaxPooling1D, Flatten, and Dense layers.
- The second CNN model (model2) has a similar architecture but with different hyperparameters such as filter size, kernel size, strides, and pool size.

#### Model Compilation and Training:



- Both models are compiled using the Adam optimizer and sparse categorical cross-entropy loss function.
- The models are trained using the training data and validated using a portion of the training data.
- The training progress is monitored over multiple epochs.

#### Model Evaluation:

- The trained models are evaluated using the test dataset to measure their performance.
- Evaluation metrics such as accuracy, loss, precision, recall, and F1-score are calculated.
- Confusion matrices are generated to visualize the model's performance in classifying different categories.

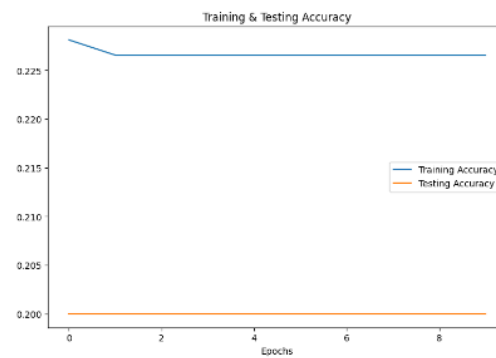
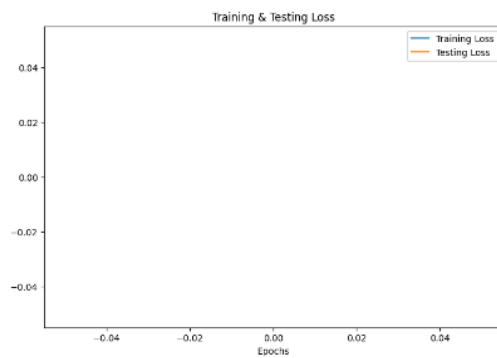
#### Results of Model 1 (128 Filters, 3 kernels, 1 Stride, Max Pool Size = 2):

- Test accuracy: 0.9400



	precision	recall	f1-score	support
1	0.91	0.99	0.95	74
2	0.97	0.92	0.95	39
3	0.96	0.96	0.96	47
4	0.94	0.85	0.89	40
accuracy			0.94	200
macro avg	0.95	0.93	0.94	200
weighted avg	0.94	0.94	0.94	200

## Results of Model 2



	precision	recall	f1-score	support
0	0.00	0.00	0.00	0.0
1	0.00	0.00	0.00	74.0
2	0.00	0.00	0.00	39.0
3	0.00	0.00	0.00	47.0
4	0.00	0.00	0.00	40.0
accuracy			0.00	200.0
macro avg	0.00	0.00	0.00	200.0
weighted avg	0.00	0.00	0.00	200.0

## Conclusion:

Model 1 performed well as compared to Model 2. As the parameters change, the metrics also varies